

graph_home_hunter_demo

June 24, 2025

1 Graph-Stop Home Hunter — Guided Notebook Tour

This notebook is a **hands-on companion** to the codebase (`graph_home_hunter.py`). You'll train, evaluate, and *peek inside* a three-stage pipeline that turns raw real-estate listings into **actionable buy/wait advice**.

Stage	What you'll do	Why it matters
1 Graph Encoder	Train a 2-layer GAT (or swap in GraphSAGE) on a buyer → listing → suburb graph.	Captures relational cues—price clusters, suburb similarity, CBD distance—that flat features miss.
2 Tabular Head	Feed embeddings + one-hot features through a Gaussian NB or Logistic Reg head, optionally keeping only the top-k features by mutual information.	Blends learned context with interpretable/calibrated probabilities.
3 Decision Logic	Apply a threshold and a Bayesian Option-Value Planner (wait-cost vs regret-cost) to decide “ <i>Bid now</i> ” vs “ <i>Hold for better stock.</i> ”	Converts probabilities into real decisions, balancing confidence against future opportunity.

1.0.1 Notebook flow

1. **Synthetic data** Generate or load 10 weekly CSV snapshots **plus** a larger train/test table.
2. **Train & validate** • Graph-only GNN • Stacked head with calibration.
3. **Threshold sweep** Find the F1-optimal (or F- for a precision/recall tilt).
4. **Week-by-week simulation** Replay snapshots and watch the Option-Value Planner decide when to pull the trigger.
5. **Explainability** Inspect per-listing feature contributions to understand each recommendation.

Ready? Let's hunt some homes

```
[23]: import os, sys, math, random, copy, json, itertools, types, warnings, joblib, \
      ↪ subprocess
```

```

from pathlib import Path

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import matplotlib.patches as mpatches

import torch, torch.nn as nn
from torch_geometric.data import Data
from torch_geometric.loader import NeighborLoader
from torch_geometric.nn import GATConv

from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report, confusion_matrix
from sklearn.naive_bayes import GaussianNB
from sklearn.linear_model import LogisticRegression
from sklearn.calibration import CalibratedClassifierCV
from sklearn.feature_selection import mutual_info_classif

# local source (assumes notebook lives next to graph_home_hunter.py)
from graph_home_hunter import (
    ListingGraphTG, GCN, IntentHead, GraphTrainer,
    Metrics, _safe_torch_load, set_seed
)

# put this once, near the top of the notebook
import warnings
from sklearn.exceptions import ConvergenceWarning

warnings.filterwarnings("ignore", category=ConvergenceWarning)

set_seed(42)
print(torch.__version__)

```

2.3.0+cu118

1.1 1. Prepare synthetic training & test CSVs

```

[2]: # If the CSVs don't exist yet, generate them with make_training_data_v3.py
if not Path('train_listings.csv').exists():
    !python make_training_data_v3.py

df_train = pd.read_csv('train_listings.csv')
df_test  = pd.read_csv('test_listings.csv')
print(df_train.shape, df_test.shape)
df_train.head()

```

(16000, 19) (4000, 19)

```
[2]:
```

	id	asking_price	body_corp	internal_m2	balcony_m2	beds	suburb \
0	10651	634000	1809	74	1.0	2	Footscray
1	2042	557000	1894	86	17.7	2	Sunbury
2	8669	601000	1406	109	17.7	3	Werribee
3	1115	616000	2517	116	8.2	2	Footscray
4	13903	639000	2303	84	6.6	2	Richmond

	livability	condition	has_cladding	travel_time_mins	sold_date \
0	10	new	False	5	NaN
1	9	good	False	61	NaN
2	8	good	False	45	NaN
3	10	good	False	14	NaN
4	16	new	False	21	NaN

	has_parking	has_storage	has_solar	north_facing	near_shops	outdoor_space \
0	NaN	False	False	False	False	False
1	False	True	False	False	True	False
2	False	NaN	NaN	False	True	True
3	True	False	False	True	False	True
4	False	False	False	NaN	True	False

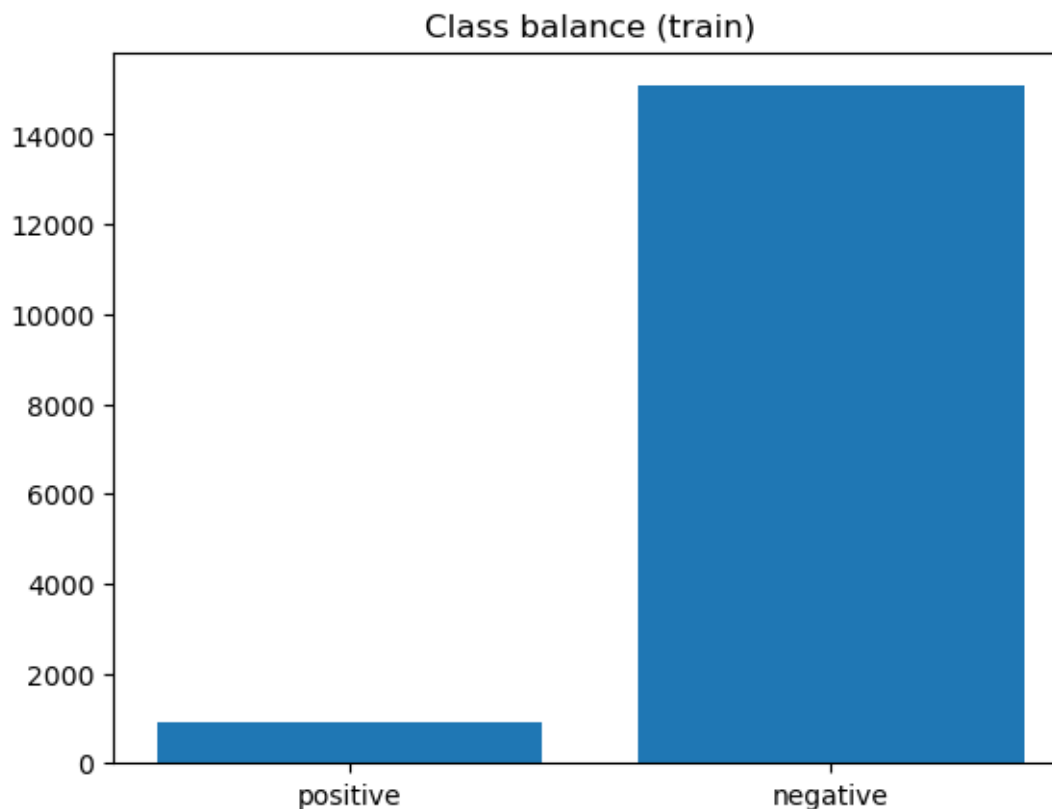
	label
0	0
1	0
2	0
3	0
4	0

1.2 2. Build a PyG graph from the training CSV

```
[3]: g_train = ListingGraphTG().from_csv('train_listings.csv')
pos = int(g_train.y[g_train.listing_idx].sum())
neg = len(g_train.listing_idx) - pos
print(f'Positives: {pos}    Negatives: {neg}')

plt.bar(['positive', 'negative'], [pos, neg])
plt.title('Class balance (train)')
plt.show()
```

Positives: 920 Negatives: 15080



1.3 3. Train the GNN encoder

```
[4]: ckpt = Path("gnn_v1.pt")

if ckpt.exists():
    print(" found", ckpt, "- loading checkpoint ...")
    gnn = GCN(); gnn.load(ckpt); gnn.eval()
else:
    print(" checkpoint not found - training from scratch ...")
    !python graph_home_hunter.py --train-graph train_graph.pt gnn_v1.pt
    ↪--epochs 150 --patience 50 --lr 1e-3
    gnn = GCN(); gnn.load(ckpt); gnn.eval()
```

found gnn_v1.pt - loading checkpoint ...

1.4 4. Pure-GNN performance on the held-out test set

```
[5]: g_test = ListingGraphTG().from_csv('test_listings.csv', add_masks=False)
with torch.no_grad():
    probb_gnn = torch.sigmoid(gnn(g_test)).numpy()[g_test.listing_idx]
    pred_gnn = (probb_gnn>=0.5).astype(int)
```

```
Metrics.report('GNN only', g_test.y[g_test.listing_idx].int().numpy(), pred_gnn)
```

GNN only acc=0.953 prec=0.57 rec=0.87 f1=0.69
confusion: TP 207 | FP 157 | FN 30 | TN 3606

1.5 5. Train tabular head on top of frozen GNN embeddings

```
[6]: head_ckpt = Path("logreg_head.pkl")

if head_ckpt.exists():
    print(" found", head_ckpt, "- loading calibrated head ...")
else:
    print(" head checkpoint not found - training now ...")
    !python graph_home_hunter.py --train-stack train_listings.csv gnn_v1.pt
    ↪ logreg_head.pkl --stack-model logreg --calib sigmoid
```

```
head = IntentHead.load(head_ckpt)
```

found logreg_head.pkl - loading calibrated head ...

1.6 6. Stacked model performance

```
[7]: !python graph_home_hunter.py --predict-stack test_listings.csv gnn_v1.pt
    ↪ logreg_head.pkl
```

stack metrics:
stack acc=0.991 prec=0.94 rec=0.90 f1=0.92
confusion: TP 213 | FP 14 | FN 24 | TN 3749

OVP decision: **PLACE OFFER NOW**

- current best posterior : 0.999
- P(better next week) : 0.68%
- Δ (Expected regret) : -0.0193

1.7 7. Model architecture schematic

```
[21]: labels = [
    "Listing graph\n(GNN input)",
    "2-layer GAT\nencoder",
    "Node\nembeddings",
    "Concat\n(+ one-hot tabular)",
    "Head classifier\n(LogReg / G-NB)",
    "Calibrator\n(sigmoid / isotonic)",
    "Option-Value\nPlanner (OVP)"
]

# styling -----
FACE, EDGE = "#d7e8ff", "#2074ff"
```

```

fs_title, fs_lbl = 16, 10.5
h_pad, v_pad = 0.04, 0.24          # figure-side padding
gap = 0.035                        # gap between blocks
arrow_kw      = dict(arrowstyle="->", lw=2.2, color="black")
arrow_final_kw = dict(arrowstyle="->", lw=2.2, color="crimson")

# geometry -----
n = len(labels)
avail_w = 1 - 2 * h_pad
box_w    = (avail_w - gap * (n - 1)) / n
box_h    = 1 - 2 * v_pad
y0       = v_pad                    # bottom of blocks

fig, ax = plt.subplots(figsize=(16, 3.8))
ax.axis("off")
ax.set_title("Graph-Stop Home-Hunter - end-to-end model flow",
            fontsize=fs_title, fontweight="bold", pad=14)

centres = []

# blocks -----
for i, txt in enumerate(labels):
    x0 = h_pad + i * (box_w + gap)
    centres.append(x0 + box_w / 2)

    ax.add_patch(
        mpatches.FancyBboxPatch(
            (x0, y0), box_w, box_h,
            boxstyle="round,pad=0.02,rounding_size=0.035",
            facecolor=FACE, edgecolor=EDGE, linewidth=1.5
        )
    )
    ax.text(x0 + box_w / 2, y0 + box_h / 2, txt,
           ha="center", va="center", fontsize=fs_lbl, fontweight="bold")

# internal arrows (black) -----
for cL, cR in zip(centres[:-1], centres[1:]):
    ax.annotate("",
               xy=(cR - box_w / 2 + 0.003, 0.5),
               xytext=(cL + box_w / 2 - 0.003, 0.5),
               arrowprops=arrow_kw
    )

# final crimson arrow inside axes -----
right_edge_OVP = centres[-1] + box_w / 2
tail = (right_edge_OVP - 0.003, 0.5)
head = (min(0.99, right_edge_OVP + 0.03), 0.5) # keep head 0.99

```

```

ax.annotate("",
    xy=head, xytext=tail,
    arrowprops=arrow_final_kw,
    annotation_clip=False          # draw even if head slightly outside
)

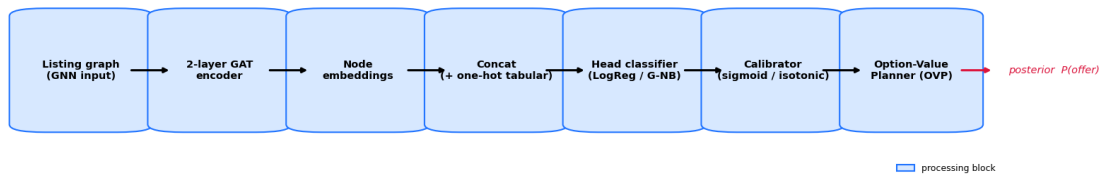
#   posterior label -----
ax.text(head[0] + 0.015, 0.5, "posterior P(offer)",
    fontsize=fs_lbl, va="center", fontstyle="italic", color="crimson")

#   legend -----
legend_patch = mpatches.FancyBboxPatch((0, 0), 1, 1,
    boxstyle="round,pad=0.02",
    facecolor=FACE, edgecolor=EDGE,
    linewidth=1.5)
ax.legend([legend_patch], ["processing block"],
    loc="lower right", frameon=False, fontsize=9)

plt.tight_layout()
plt.show()

```

Graph-Stop Home-Hunter - end-to-end model flow



1.8 8. Week-by-week simulation

```

[9]: !python simulate_weeks.py --gnn gnn_v1.pt --head logreg_head.pkl --thr 0.95
    ↪ --wait-cost 0.02 --regret-cost 0.2 week_0.csv week_1.csv week_2.csv week_3.
    ↪ csv week_4.csv week_5.csv week_6.csv week_7.csv week_8.csv week_9.csv

```

```

Week 0: c:\Users\User\anaconda3\python.exe graph_home_hunter.py --predict-
stack week_0.csv gnn_v1.pt logreg_head.pkl --thr 0.95 --wait-cost 0.02 --regret-
cost 0.2 --k-next-week 10 --quiet-metrics

```

```

OVP decision:      WAIT
• current best posterior : 0.747
• P(better next week)   : 17.13%
• Δ(Expected regret)    : +0.0143

```

Week 1: c:\Users\User\anaconda3\python.exe graph_home_hunter.py --predict-stack week_1.csv gnn_v1.pt logreg_head.pkl --thr 0.95 --wait-cost 0.02 --regret-cost 0.2 --k-next-week 10 --quiet-metrics

OVP decision: WAIT

- current best posterior : 0.747
- P(better next week) : 17.12%
- Δ (Expected regret) : +0.0142

Week 2: c:\Users\User\anaconda3\python.exe graph_home_hunter.py --predict-stack week_2.csv gnn_v1.pt logreg_head.pkl --thr 0.95 --wait-cost 0.02 --regret-cost 0.2 --k-next-week 10 --quiet-metrics

OVP decision: WAIT

- current best posterior : 0.747
- P(better next week) : 17.11%
- Δ (Expected regret) : +0.0142

Week 3: c:\Users\User\anaconda3\python.exe graph_home_hunter.py --predict-stack week_3.csv gnn_v1.pt logreg_head.pkl --thr 0.95 --wait-cost 0.02 --regret-cost 0.2 --k-next-week 10 --quiet-metrics

OVP decision: WAIT

- current best posterior : 0.747
- P(better next week) : 16.96%
- Δ (Expected regret) : +0.0139

Week 4: c:\Users\User\anaconda3\python.exe graph_home_hunter.py --predict-stack week_4.csv gnn_v1.pt logreg_head.pkl --thr 0.95 --wait-cost 0.02 --regret-cost 0.2 --k-next-week 10 --quiet-metrics

OVP decision: WAIT

- current best posterior : 0.570
- P(better next week) : 20.45%
- Δ (Expected regret) : +0.0209

Week 5: c:\Users\User\anaconda3\python.exe graph_home_hunter.py --predict-stack week_5.csv gnn_v1.pt logreg_head.pkl --thr 0.95 --wait-cost 0.02 --regret-cost 0.2 --k-next-week 10 --quiet-metrics

OVP decision: WAIT

- current best posterior : 0.570

- P(better next week) : 20.43%
- Δ (Expected regret) : +0.0209

Week 6: c:\Users\User\anaconda3\python.exe graph_home_hunter.py --predict-stack week_6.csv gnn_v1.pt logreg_head.pkl --thr 0.95 --wait-cost 0.02 --regret-cost 0.2 --k-next-week 10 --quiet-metrics

OVP decision: WAIT

- current best posterior : 0.934
- P(better next week) : 10.94%
- Δ (Expected regret) : +0.0019

Week 7: c:\Users\User\anaconda3\python.exe graph_home_hunter.py --predict-stack week_7.csv gnn_v1.pt logreg_head.pkl --thr 0.95 --wait-cost 0.02 --regret-cost 0.2 --k-next-week 10 --quiet-metrics

OVP decision: WAIT

- current best posterior : 0.946
- P(better next week) : 10.47%
- Δ (Expected regret) : +0.0009

Week 8: c:\Users\User\anaconda3\python.exe graph_home_hunter.py --predict-stack week_8.csv gnn_v1.pt logreg_head.pkl --thr 0.95 --wait-cost 0.02 --regret-cost 0.2 --k-next-week 10 --quiet-metrics

OVP decision: **PLACE OFFER NOW**

- current best posterior : 0.954
- P(better next week) : 9.96%
- Δ (Expected regret) : -0.0001

Week 9: c:\Users\User\anaconda3\python.exe graph_home_hunter.py --predict-stack week_9.csv gnn_v1.pt logreg_head.pkl --thr 0.95 --wait-cost 0.02 --regret-cost 0.2 --k-next-week 10 --quiet-metrics

OVP decision: WAIT

- current best posterior : 0.566
- P(better next week) : 21.80%
- Δ (Expected regret) : +0.0236

1.9 9. Explain why the best listing was chosen in Week 8

```
[10]: !python graph_home_hunter.py --predict-stack week_8.csv gnn_v1.pt logreg_head.  
      ↪pkl --thr 0.95 --explain-id 161
```

```
stack metrics:  
stack acc=0.972  prec=1.00  rec=0.17  f1=0.29  
confusion: TP 1 | FP 0 | FN 5 | TN 175
```

```
OVP decision:      WAIT  
• current best posterior : 0.954  
• P(better next week)   : 26.94%  
• Δ(Expected regret)    : +0.0069
```

```
Top feature contributions for listing 161:  
near_shops_True          +4.516  
condition_new            +3.104  
north_facing_True       +0.343  
outdoor_space_True      -0.312  
has_storage_False       -0.236
```

1.10 10. Run the Streamlit GUI (opens a browser tab)

1.10.1 To run the interactive simulations in the GUI, please upload:

- 1) The model: gnn_v1.pt
- 2) The model head: logreg_head.pkl and
- 3) The mock weekly listings data: week0-9.csv files

```
[11]: !streamlit run home_hunter_app.py
```

^C